

# Planning

# Outline

- Introduction to the planning problem
- Formulating planning problems
- Languages for planning problems
- Planning algorithm - search in plan space

# Planning

- The task of coming up with a sequence of actions that will achieve a goal
- What plans do you have for this evening/this weekend?
- Why should an intelligent agent be able to plan?
  - If you fail to plan, you plan to fail  
Benjamin Franklin
- Agents with planning capabilities differ from pure reactive agents
  - Which type is better?

# Applications of Planning

- Mobile Robots
- Virtual agents
- Action choice + resource handling
  - for transportation of goods
  - at schools, hospitals
  - Hubble Space Telescope scheduler
  - Work-flow management
- Software test case generation

# Plans and Actions in Jason Agent Programming Language

```
/* Initial beliefs and rules */
current_tactic(give_and_go(su_monday, ras_ruby)).

/* Initial goals */
!start_playing.

/* Plans */

+!start_playing <- connect_to_SL("rasika", "Wonderland NW,178,164,1001");
                    !check_connected.

+!check_connected: not connected <-wait(2000);
                    !!check_connected.

+!check_connected: connected <- action("run","210,241,1001").

+successful_pass_by_upkick(OtherAgent, Me)[state(N)]:current_tactic(give_and_go(OtherAgent, Me))& .my_name(Me)
  <- .term2string(OtherAgent,OtherAgentStr);
     .concat(OtherAgentStr, " passed ME the ball by an UP KICK: start monitoring for expectation", Msg);
     msg_sl(Msg);
     .start_monitoring("fulf", "reach_penaltyB", "expectation_monitor", NegShootRangeConjunction, FulfProposition, [N])

+fulf("fulf_reach_penaltyB", N) <-msg_sl("Expectation FULFILLED: pass the ball back");
                                .stop_monitoring("fulf", "reach_penaltyB", "expectation_monitor");
                                .stop_monitoring("viol", "reach_penaltyB", "expectation_monitor");
                                action("pass","up shot").

+viol("viol_reach_penaltyB", N)
  <- .stop_monitoring("fulf", "reach_penaltyB", "expectation_monitor");
     .stop_monitoring("viol", "reach_penaltyB", "expectation_monitor");
     msg_sl("Expectation VIOLATED: run to penalty area");
     !choose_and_enact_new_tactic.

+!choose_and_enact_new_tactic : .my_name(Me) <-
  +current_tactic(solo_run(Me));
  action("run","54,101,790");
  +current_tactic(give_and_go(su_monday, ras_ruby)).
```

# Planning

- Given the current **state** and the state of the world you want to achieve (**goals**),
  - determine how (and when) to do it (**plan**).
  - I.e. Figuring out a sequence of **actions** to achieve the goal
- We will restrict this discussion to **deterministic**, **fully observable**, **static** and **discrete situations** => Classical Planning
  - Static => Changes occur only when the agent acts

# A Planning Problem

- Goal: Have a birthday party
- Current state:
  - Agent is at home
  - There is enough flour in the larder
  - But no butter or sugar
- Things to do
  - Invite friends
  - Make a cake

# Planning Problem

- Given:
  - A way to describe the world. I.e. state representation
  - An initial state of the world
  - A goal description
  - A set of possible actions to change the world
- Find:
  - A sequence of actions to go from initial state to a state where the goal is achieved

Planning problem



# Planning Requirements

- A way to express the planning problem
- An algorithm to solve the problem
- Planning algorithms based on logical approaches should take advantage of the logical structure of the problem
  - For this, the problem should be expressed in a suitable logical language

# Language of Planning Problems

- The language should be
  - Expressive enough to describe a variety of problems
  - Restrictive enough to let efficient algorithms to operate on it
- It should represent **states**, **goals** and **actions**
- Basic language of classical planners - PDDL (Planning Domain Definition Language)
-

# PDDL Syntax

- **State** representation
  - Conjunction of **ground atomic fluents**
  - Literals can be
    - Propositional ( $Poor \wedge Unknown$ )
    - First order ( $At(Plane_1, Sydney) \wedge At(Plane_2, Perth)$ )
  - First order literals must be
    - Ground ( $At(x,y)$ ) ❌
    - Function-free ( $At(Father(Fred), Perth)$ ) ❌
  - **Closed-world assumption:** Whatever is not mentioned is assumed to be false
  - Unique name assumption:  $Plane_1 \neq Plane_2$

# PDDL Syntax

- **Goal** representation
  - Partially specified state
    - E.g.
      - *At (Plane<sub>1</sub>, Melbourne)*
      - *Rich  $\wedge$  Famous*
  - A propositional state  $s$  satisfies goal  $g$  if  $s$  contains all the atoms of  $g$
  - E.g. State *(Rich  $\wedge$  Famous  $\wedge$  Happy)* satisfies the goal *(Rich  $\wedge$  Famous)*

# PDDL Syntax

- **Action** representation
  - Specified using an action schema containing
    - **Preconditions**: must hold to execute the action
    - **Effects**: changes caused by the action
  - E.g.

Action ( *Fly*( *p*, *from*, *to*),

PRECOND:  $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT:  $\neg At(p, from) \wedge At(p, to)$

# PDDL Syntax

- Action name and parameter list
- Identifies the action

Action ( *Fly*( *p*, *from*, *to*),

PRECOND: *At* (*p*, *from*)  $\wedge$  *Plane*(*p*)  $\wedge$   
*Airport*(*from*)  $\wedge$  *Airport*(*to*)

EFFECT:  $\neg$ *At* (*p*, *from*)  $\wedge$  *At* (*p*, *to*)

- States what should be true in a state for the action to execute
- Conjunction of function-free (+)ve literals
- Variables should appear in the parameter list

- Describes the state change when the action is executed
- Contains function-free literals
- Variables should appear in the parameter list
- Can be broken into an **add list** and a **delete list**

Action ( *Fly*( *p*, *from*, *to*),

PRECOND:  $At(p, from) \wedge Plane(p) \wedge Airport(fr)$   
 $\wedge Airport(to)$

EFFECT:  $\neg At(p, from) \wedge At(p, to)$

Action ( *Fly*( *p*, *from*, *to*),  
PRECOND:  $At(p, from) \wedge Plane(p) \wedge$   
Airport(*from*)  $\wedge \neg Airport(to)$   
EFFECT:  $\neg At(p, from) \wedge At(p, to)$



# PDDL Semantics

- An action is **applicable** in any state that satisfies the precondition
- How to establish applicability?

**State** -  $At(P_1, JFK) \wedge At(P_2, SFO) \wedge Plane(P_1) \wedge$   
 $Plane(P_2) \wedge Airport(JFK) \wedge Airport(SFO)$

**satisfies**

**Precondition** -  $At(p, from) \wedge Plane(p) \wedge Airport(from)$   
 $\wedge Airport(to)$

**If substituted with**

$\{p/P_1, from/JFK, to/SFO\}$

# PDDL Semantics

- Starting at state  $s$ , **result** of executing action  $a$  is a state  $s'$ 
  - Contains (+)ve literals in the effect of  $a$
  - (-)ve literals are removed
  - Otherwise same as  $s$
- **Assumption** – every literal not mentioned in the effect remains unchanged
  - Avoids **frame problem**
- Solution to the planning problem – an action sequence that, when executed, results in a state that satisfies the goal

# ADL (Action Description Language)

- In recent years, it has been noted that STRIPS is *not* sufficiently expressive
- This has resulted in many language variants. ADL is one of them.
- The same *Fly* action in ADL

Action(*Fly*(*p:Plane*, *from:Airport*, *to:Airport*),

PRECOND:  $At(p, from) \wedge (from \neq to)$

EFFECT:  $\neg At(p, from) \wedge At(p, to)$

# STRIPS Vs ADL

Action ( *Fly*( *p*, *from*, *to*),

PRECOND:  $At(p, from) \wedge Plane(p) \wedge$   
 $Airport(from) \wedge Airport(to)$

EFFECT:  $\neg At(p, from) \wedge At(p, to)$

Action(*Fly*(*p:Plane*, *from:Airport*, *to:Airport*),

– PRECOND:  $At(p, from) \wedge (from \neq to)$

– EFFECT:  $\neg At(p, from) \wedge At(p, to)$

# STRIPS Vs ADL

STRIPS Language	ADL Language
Only positive literals in states: <i>Poor</i> $\wedge$ <i>Unknown</i>	Positive and negative literals in states: $\neg$ <i>Rich</i> $\wedge$ $\neg$ <i>Famous</i>
Closed World Assumption: Unmentioned literals are false.	Open World Assumption: Unmentioned literals are unknown.
Effect $P \wedge \neg Q$ means add $P$ and delete $Q$ .	Effect $P \wedge \neg Q$ means add $P$ and $\neg Q$ and delete $\neg P$ and $Q$ .
Only ground literals in goals: <i>Rich</i> $\wedge$ <i>Famous</i>	Quantified variables in goals: $\exists x \text{ At}(P_1, x) \wedge \text{At}(P_2, x)$ is the goal of having $P_1$ and $P_2$ in the same place.
Goals are conjunctions: <i>Rich</i> $\wedge$ <i>Famous</i>	Goals allow conjunction and disjunction: $\neg$ <i>Poor</i> $\wedge$ ( <i>Famous</i> $\vee$ <i>Smart</i> )
Effects are conjunctions.	Conditional effects allowed: <b>when</b> $P$ : $E$ means $E$ is an effect only if $P$ is satisfied.
No support for equality.	Equality predicate ( $x = y$ ) is built in.
No support for types.	Variables can have types, as in ( $p$ : <i>Plane</i> ).
<b>Figure 11.1</b> Comparison of STRIPS and ADL languages for representing planning problems. In both cases, goals behave as the preconditions of an action with no parameters.	

# Example Planning Problem: Air Cargo

$Init(At(C_1, SFO) \wedge At(C_2, JFK) \wedge At(P_1, SFO) \wedge At(P_2, JFK)$   
 $\wedge Cargo(C_1) \wedge Cargo(C_2) \wedge Plane(P_1) \wedge Plane(P_2)$   
 $\wedge Airport(JFK) \wedge Airport(SFO))$

$Goal(At(C_1, JFK) \wedge At(C_2, SFO))$

$Action(Load(c, p, a),$

PRECOND:  $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT:  $\neg At(c, a) \wedge In(c, p)$ )

$Action(Unload(c, p, a),$

PRECOND:  $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT:  $At(c, a) \wedge \neg In(c, p)$ )

$Action(Fly(p, from, to),$

PRECOND:  $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT:  $\neg At(p, from) \wedge At(p, to)$ )

$[Load(C_1, P_1, SFO), Fly(P_1, SFO, JFK), Unload(C_1, P_1, JFK),$   
 $Load(C_2, P_2, JFK), Fly(P_2, JFK, SFO), Unload(C_2, P_2, SFO)]$

# Example Planning Problem: The Spare Tire Problem

*Init*(*Tire*(*Flat*)  $\wedge$  *Tire*(*Spare*)  $\wedge$  *At*(*Flat*, *Axle*)  $\wedge$  *At*(*Spare*, *Trunk*))  
*Goal*(*At*(*Spare*, *Axle*))  
*Action*(*Remove*(*obj*, *loc*),  
    PRECOND: *At*(*obj*, *loc*)  
    EFFECT:  $\neg$  *At*(*obj*, *loc*)  $\wedge$  *At*(*obj*, *Ground*))  
*Action*(*PutOn*(*t*, *Axle*),  
    PRECOND: *Tire*(*t*)  $\wedge$  *At*(*t*, *Ground*)  $\wedge$   $\neg$  *At*(*Flat*, *Axle*)  
    EFFECT:  $\neg$  *At*(*t*, *Ground*)  $\wedge$  *At*(*t*, *Axle*))  
*Action*(*LeaveOvernight*,  
    PRECOND:  
    EFFECT:  $\neg$  *At*(*Spare*, *Ground*)  $\wedge$   $\neg$  *At*(*Spare*, *Axle*)  $\wedge$   $\neg$  *At*(*Spare*, *Trunk*)  
             $\wedge$   $\neg$  *At*(*Flat*, *Ground*)  $\wedge$   $\neg$  *At*(*Flat*, *Axle*)  $\wedge$   $\neg$  *At*(*Flat*, *Trunk*))

[*Remove*(*Flat*, *Axle*), *Remove*(*Spare*, *Trunk*), *PutOn*(*Spare*, *Axle*)]

# Planning Algorithms

- Use state-space search
  - Actions in a planning problem specify both preconditions and effects
  - Search can be performed in either direction
    - Forward state-space search
    - Backward state-space search



# Planning Algorithms

- State-space search
  - Forward state-space search
  - Backward state-space search
- Plan space search
  - Partial order planning (POP)

# State-space Search

- Possible because actions in a planning problem specify both preconditions and effects
- Nodes = states of the world
- Transitions between nodes = actions
- Path through the state space = plan

# Forward state-space search

- Step cost is usually one per action. But it is easy to allow differential costs
- As usual, the main issue with Forward search is that it cannot avoid irrelevant actions
  - Thus it is inefficient
  - Need to use a good heuristic function to estimate distance from state to goal

# Forward state-space search

- Also called progression planning
- Starting from the ***initial state***, consider sequences of actions that reach the goal
- For a given state, all ***actions*** whose preconditions are satisfied, are applicable
  - The successor state is determined by adding the positive effect literals and deleting the negative effect literals to the current state
- ***Goal test*** checks if the current state satisfies the goal

# Forward State-space Search

Forward-search( $O, s_0, g$ )

# Forward State-space Search

- Forward search is **sound**
  - For any plan returned by any of its nondeterministic traces, this plan is guaranteed to be a solution
- Forward search is **complete**
  - If a solution exists then at least one of Forward search's nondeterministic traces will return a solution.

# Deterministic Implementations

- Breadth-first search
- Depth-first search
- Best-first search (e.g.,  $A^*$ )
- Greedy search

# Loop Checking

- Can be done by keeping a record of the state sequence  $(s_0, s_1, \dots, s_k)$  on the current path
- Modify the algorithm to return failure if there are nodes  $i$  and  $k$  where
  - $i < k$  s.t.  $s_k = s_i$
  - $i < k$  s.t.  $s_k \subseteq s_i$



# Branching Factor


- Number of applicable plans that don't progress towards goal
- Forward search has a large branching factor
- Example:
  - Blocks world

# Blocks World

- Consists of a set of cube-shaped blocks sitting on a table.
- The blocks can be stacked
  - Only one block can fit directly on top of another
- A robot arm can pick up a block and move it to
  - The table OR
  - Top of another block
- The arm can pick up only one block at a time
  - Cannot pick up a block that has another one on it
- Goal => build one or more stacks of blocks according to the given order

# Blocks World


- Possible state predicates
  - Holding(a)
  - handEmpty
  - onTable(a)
  - on(a, b)
  - clear(a)
- Possible actions
  - Pickup(a)
  - Putdown(a)
  - Stack(a, b)
  - unstack(a, b)



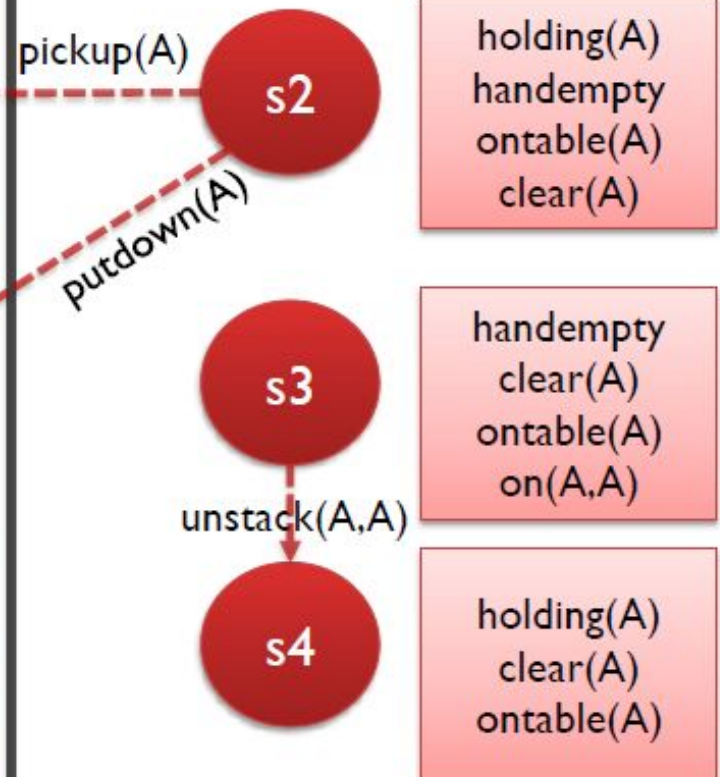
# Blocks World - Size 1

We assume we know the initial state  
We only consider states *reachable* from there!

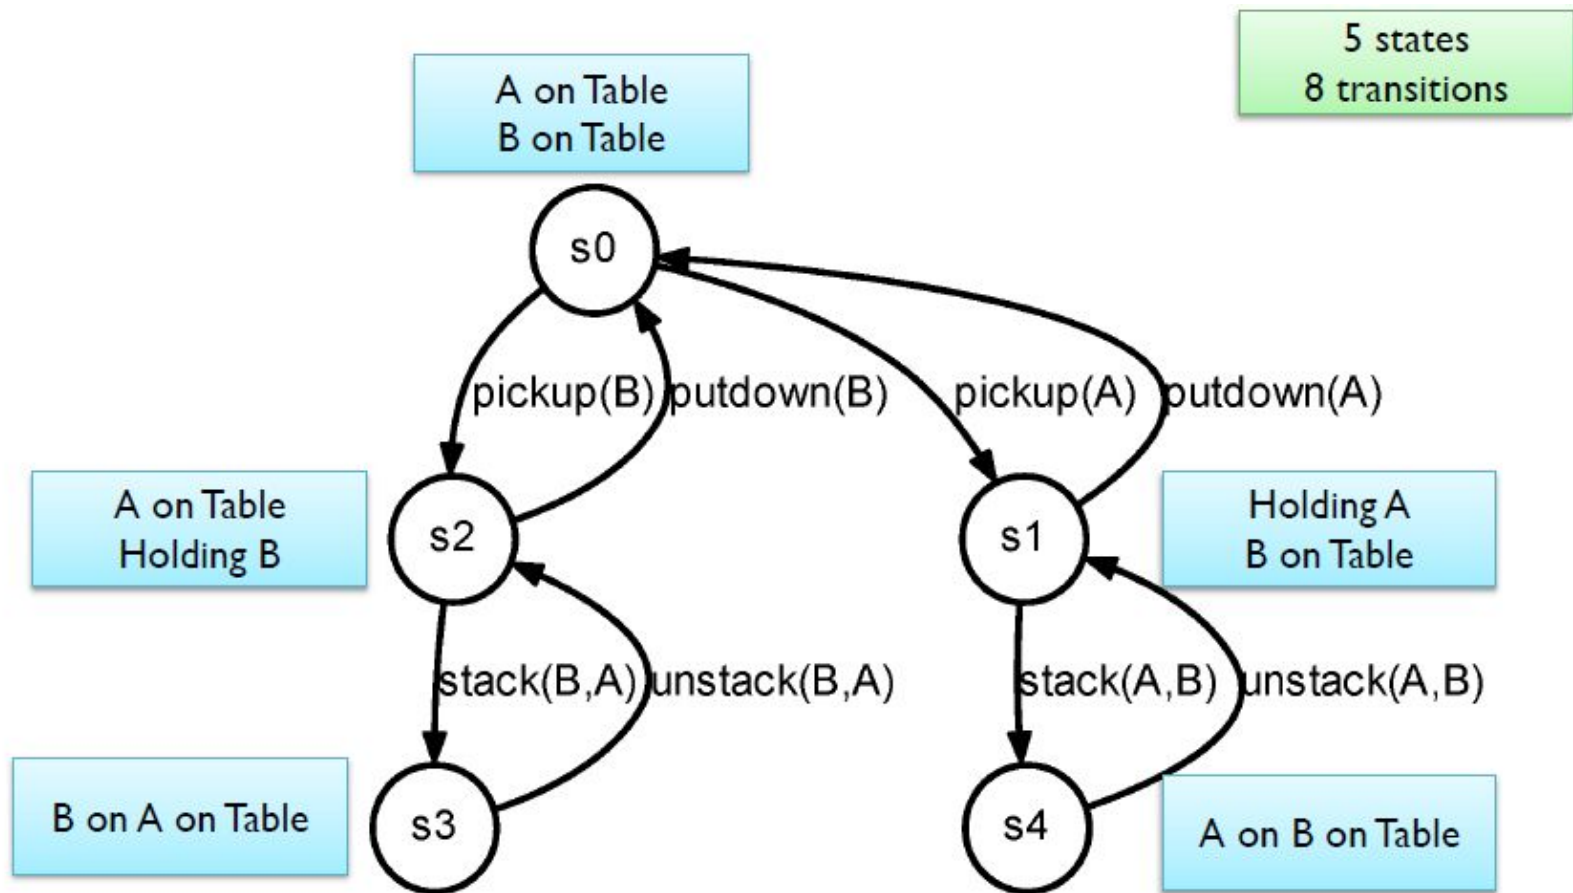
Here: Start with  $s_0$  = all blocks on the table



Many other states "exist",  
but are not reachable  
from the current starting state



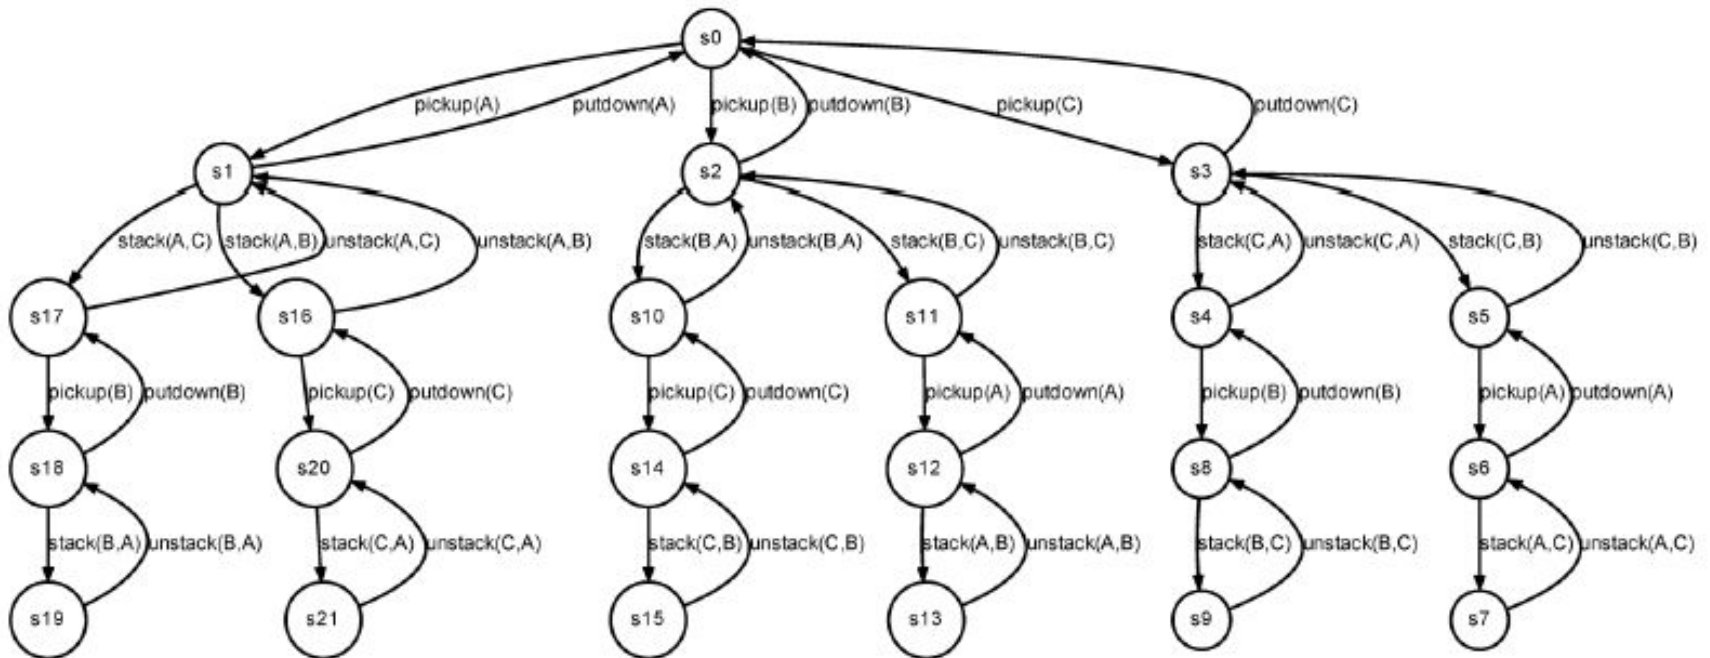
# Blocks World - Size 2



# Blocks World - Size 3


22 states  
42 transitions

A on Table  
B on Table  
C on table




# Blocks World - Size 4

125 states  
272 transitions





# Blocks World - Size 5



866 states  
2090 transitions



# Blocks World - Size 0-8

Blocks	States	States reachable from "all on table"	Transitions reachable
0	2	1	0
1	32	2	2
2	2048	5	8
3	524288	22	42
4	536870912	125	272
5	2199023255552	866	2090
6	36028797018963968	7057	18552
7	2361183241434822606848	65990	186578
8	618970019642690137449562112	695417	2094752
9	649037107316853453566312041152512	...	...
10	2722258935367507707706996859454145691648	...	...

# Forward State Space Search

- Solution??
  - Domain-specific : search control rules, heuristics
  - Domain-independent : heuristics automatically generated from the problem description

# Heuristic

A **heuristic** technique (/hjuˈrɪstɪk/; **Ancient Greek**: εὕρισκω, "find" or "discover"), often called simply a *heuristic*, is any **approach to problem solving**, learning, or discovery that **employs a practical methodology not guaranteed to be optimal or perfect, but sufficient for the immediate goals**. Where finding an optimal solution is impossible or impractical, heuristic methods can be used to speed up the process of finding a satisfactory solution. Heuristics can be mental shortcuts that ease the cognitive load of making a decision. Examples of this method include using a **rule of thumb**, an **educated guess**, an intuitive judgment, stereotyping, **profiling**, or **common sense**.

# Use of Heuristics


Example: 3 blocks, all on the table in s0



We now have  
1 *open node*,  
which is *unexpanded*

# Use of Heuristics

We visit  $s_0$  and expand it



We now have  
3 open nodes,  
which are *unexpanded*

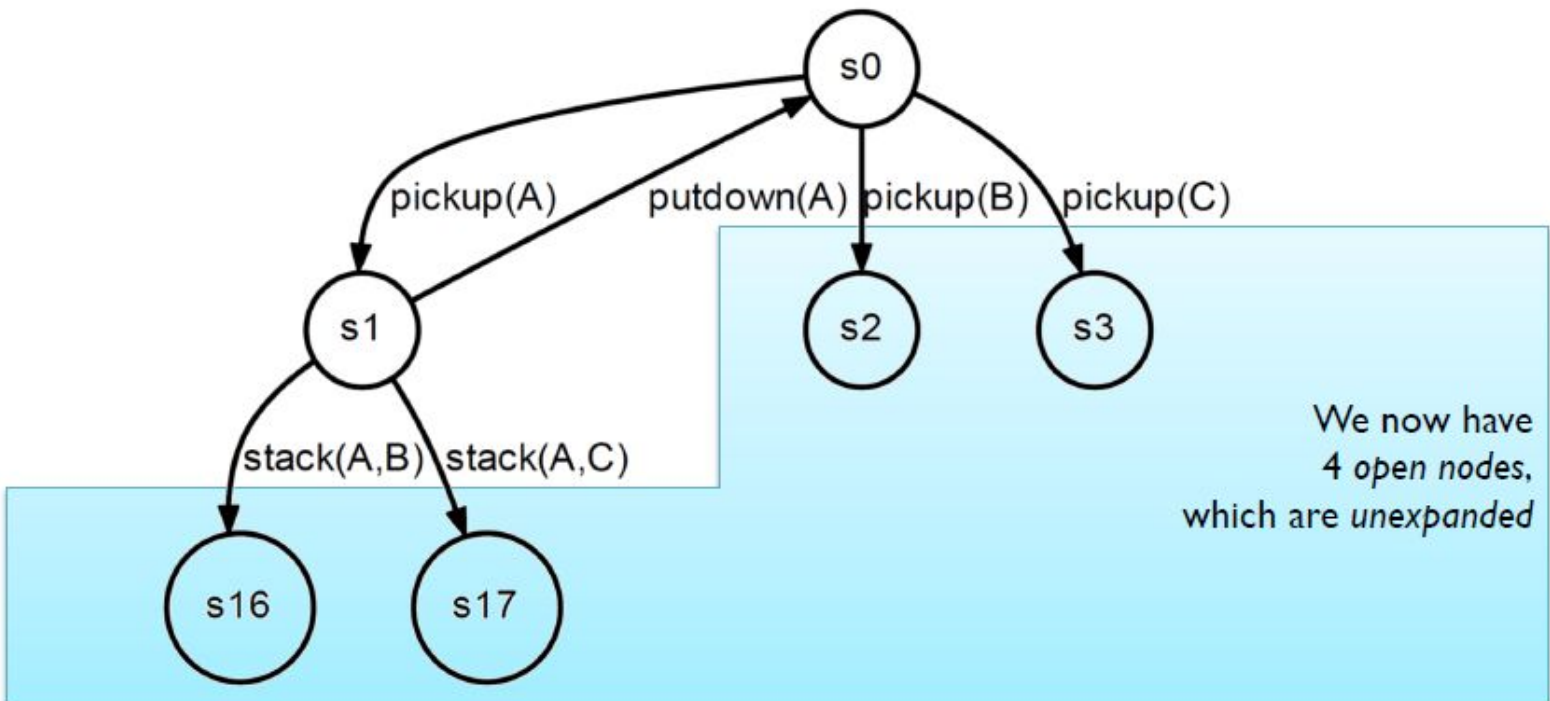
A **heuristic function** estimates the distance from each open node to the goal:

We calculate  $h(s_1)$ ,  $h(s_2)$ ,  $h(s_3)$

A **search strategy** uses this value (and other info) to *prioritize* between them

# Use of Heuristics

If we choose to visit s1:



2 new heuristic values are calculated:  $h(s16)$ ,  $h(s17)$   
The **search strategy** now has 4 nodes to prioritize

# Use of Heuristics

- Two aspects of heuristic guidance
  - Use a search strategy that can make use of heuristics
    - E.g. A\*
  - Generating the actual heuristic guidance
    - Should be able to apply on the selected search strategy
- Heuristic function should be efficient
- Two **distinct** objectives for heuristic selection
  - Find a solution quickly
  - Find a good (cheap) solution



# Use of Heuristics

## QUICK SOLUTION

Accumulated plan cost = 50  
Estimated distance to goal = 10

## GOOD (CHEAP) SOLUTION

Accumulated plan cost = 5  
Estimated distance to goal = 30



# Domain-Independent Heuristics

- Decide how close a state  $s$  is to the goal
  - Count how many facts are different
- No action costs, for simplicity
- Search strategy
  - Choose an open node with a minimal number of goal facts to achieve



# Heuristic Functions

- Which is better?
  - Domain-specific??
  - Domain-independent??

# Backward state-space search

- Also called regression planning
- Work backward from the goal(s)
- Generates only relevant actions
  - An action is relevant to a conjunctive goal, if it achieves one of the conjuncts of the goal
  - E.g. if the goal is:  $At(C_1, B) \wedge At(C_2, B) \wedge At(C_3, B)$ , then any action that satisfies  $At(C_1, B)$ ,  $At(C_2, B)$  OR  $At(C_3, B)$  is relevant to the goal.
- Actions **must** be chosen such that they are **consistent**
  - An action must *not* undo any desired literals

# Backward State Space Search

- Start at the goal and compute inverse state transitions
  - New set of subgoals =  $\gamma^{-1}(g, a)$
- To define  $\gamma^{-1}(g, a)$ , must first define relevance:
- An action  $a$  is relevant for a goal  $g$  if
  - $a$  makes at least one of  $g$ 's literals true
    - $g \cap effects(a) \neq \emptyset$
  - $a$  does not make any of  $g$ 's literals false
    - $g^+ \cap effects^-(a) = \emptyset$  and  $g^- \cap effects^+(a) = \emptyset$

# Backward State Space Search

- If  $a$  is relevant for  $g$ , then
  - $\gamma^{-1}(g, a) = (g - effects(a)) \cup precondition(a)$
- Otherwise  $\gamma^{-1}(g, a)$  is undefined

# Backward State Space Search

Backward-search( $O, s_0, g$ )

# Efficiency of Backward Searching

- Backward search can also have a very large branching factor
- As before, deterministic implementations can waste lots of time trying all of them
- solution??
  - Lifting
  - STRIPS algorithm