

# Intermediate SQL

---

CS3043 - Database Systems

# Overview

- Joins
  - Outer joins
  - Inner joins
- Views
- Transactions
- Integrity Constraints
- Referential Integrity
- Built-in Data Types
- User Defined Data Types
- Domains
- Large Object Data Types
- Indexing

# Joined Relations

- **Join operations** take two relations and return another relation as a result.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition).
- It also specifies the attributes that are present in the result of the join.
- The join operations are typically used as subquery expressions in the **from** clause.

Examples:

```
SELECT *  
FROM course LEFT OUTER JOIN prereq  
ON course.course_id = prereq.prereq_id
```

```
SELECT *  
FROM course INNER JOIN prereq  
USING (course_id)
```

# Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.

# Join operations - Example

- course relation

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- prereq relation

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that,
  - information of CS-315 is missing on prereq
  - information of CS-347 is missing on course

# LEFT OUTER JOIN


- Returns **all the rows from the left table** with matching rows from right table. If there is no match in the right table, those values will be null.

Examples:

```
SELECT *  
FROM course LEFT OUTER JOIN prereq  
ON course.course_id = prereq.course_id
```

```
SELECT *  
FROM course LEFT OUTER JOIN prereq  
USING (course_id)
```

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null



# RIGHT OUTER JOIN


- Returns **all the rows from the right table** with matching rows from left table. If there is no match in the left table, those values will be null.

Examples:

```
SELECT *  
FROM course RIGHT OUTER JOIN prereq  
ON course.course_id = prereq.course_id
```

```
SELECT *  
FROM course RIGHT OUTER JOIN prereq  
USING (course_id)
```

<i>course_id</i>	<i>prereq_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	BIO-101	Genetics	Biology	4
CS-190	CS-101	Game Design	Comp. Sci.	4
CS-347	CS-101	<i>null</i>	<i>null</i>	<i>null</i>




# FULL OUTER JOIN

- Returns **all the rows from the right table and all the rows from the left table**. Missing values will be null.

Examples:

```
SELECT *  
FROM course FULL OUTER JOIN prereq  
ON course.course_id = prereq.course_id
```



- Not supported in MySQL

How do you emulate FULL OUTER JOIN in MySQL using LEFT and RIGHT JOIN?




# INNER JOIN

- Returns the rows when there is **a match in both tables**.

Examples:

```
SELECT *  
FROM course INNER JOIN prereq  
ON course.course_id = prereq.course_id
```

```
SELECT *  
FROM course INNER JOIN prereq  
USING (course_id)
```



# Writing some queries

*Instructor(ID, name, dept\_name, salary)*

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

*Teaches(ID, course\_id, sec\_id, semester, year)*

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

Using JOINS only,

1. Find the names of instructors who do not teach any courses

**select** name

**from** instructor **left outer join** teaches **using**(ID)

**where** course\_id **is null**;

2. Find the names of instructors who teach at least one course (no duplicates)

**select distinct** name

**from** instructor **inner join** teaches **using**(ID);

# Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations) stored in the database.
- Consider a person who needs to know an instructor's name and department, but not the salary. This person should see a relation described in SQL by,

```
SELECT ID, name, dept_name  
FROM instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.

# VIEW definition

- A view is defined using the **create view** statement which has the form

**CREATE VIEW** *v* **AS** < query expression >

*v* - name of the view

<query expression> - any legal SQL expression

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
  - Rather, a view definition causes the saving of an expression; **the expression is substituted into queries using the view.**

# Example Views

- A view of instructors without their salary

```
create view faculty as  
select ID, name, dept_name _name  
from instructor
```

- Find all instructors in the Biology department

```
select name  
from faculty  
where dept_name = 'Biology'
```

- Create a view of department salary totals

```
create view departments_total_salary(dept_name, total_salary) as  
select dept_name, sum(salary)  
from instructor  
group by dept_name;
```

# View dependencies

- One view may be used in the expression defining another view.
- A view relation  $v_2$  is said to **depend directly** on a view relation  $v_1$ , if  $v_1$  is used in the expression defining  $v_2$

$$V_1 \leftarrow V_2$$

- A view relation  $v_2$  is said to **depend on** view relation  $v_1$  if either  $v_2$  depends directly on  $v_1$  or there is a path of dependencies from  $v_2$  to  $v_1$

$$V_1 \leftarrow V_2 \quad \text{or} \quad V_1 \leftarrow \dots \leftarrow V_2$$

- A view relation  $v$  is said to be **recursive** if it depends on itself.

# Views Defined Using Other Views

```
CREATE VIEW physics_fall_2009 AS  
  SELECT course.course_id, sec_id, building, room_number  
  FROM course, section  
  WHERE course.course_id = section.course_id  
        AND course.dept_name = 'Physics'  
        AND section.semester = 'Fall'  
        AND section.year = '2009'
```

```
CREATE VIEW physics_fall_2009_watson AS  
  SELECT course_id, room_number  
  FROM physics_fall_2009  
  WHERE building = 'Watson'
```

# View expansion

- A way to define the meaning of views defined in terms of other views.
- Let the view  $v_1$  be defined by an expression  $e_1$  that may itself contains uses of view relations.
- View expansion of an expression repeats the following replacement step:

**repeat**

Find any view relation  $v_i$  in  $e_1$

Replace the view relation  $v_i$  by the expression defining  $v_i$

**until** no more view relations are present in  $e_1$

- As long as the view definitions are not recursive, this loop will terminate.



# Updating a View

- For a view to be updatable, there must be a one-to-one relationship between the rows in the view and the rows in the underlying table.
- Most SQL implementations allow updates only on simple views
  - The **from** clause has only one database relation.
  - The **select** clause contains only attribute names of the relation and does not have any **expressions**, **aggregates**, or **distinct** specification.
  - Any attribute not listed in the **select** clause can be set to null
  - The query does not have a **group by** or **having** clause.

Example:

```
INSERT INTO faculty VALUES (30765, 'Green', 'Music');
```

⇒ Query OK

```
INSERT INTO dept_total_salary VALUES ('Nuclear', 299000.00);
```

⇒ ERROR : The target table dept\_total\_salary of the INSERT is not insertable-into

# Materialized Views

- A logical view of the data driven by a SELECT query. But creates a physical relation/table containing all the tuples in the result of the query defining the view.
- If relations used in the query are updated, the materialized view (MV) result becomes out of date
  - Need to **maintain** the view, by updating the view whenever the underlying relations are updated.
- **Performance is higher** than Views
  - Instead of running the expanded query against the database every time, MV acquire the results from a physical table
- Not supported in MySQL

# Transactions

- Transaction is a **Unit of work** that is performed against a database.
- Have following four properties
  - **Atomicity** - the whole sequence of actions within a transaction is either fully executed or, in case of any exception, rolled back as if it never occurred.
  - **Consistency** - ensures that the database properly changes states upon a successfully committed transaction.
  - **Isolation** - enables transactions to occur independently and transparent of each other.
  - **Durability** - ensures that the result or effect of a committed transaction persists in case of a system failure.
- By default on most databases each SQL statement commits automatically
  - Can turn off auto commit for a session

**START TRANSACTION;**

<sequence of sql actions>

**COMMIT;**

# Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
- Example constraints:
  - A checking account must have a balance greater than \$10,000.00
  - A salary of a bank employee must be at least \$40 an hour
  - A customer must have a (non-null) phone number

# Integrity Constraints

- **NOT NULL**

- declare certain attributes to be not null

Example:

```
name varchar(20) NOT NULL,  
budget numeric(12,2) NOT NULL,
```

- **PRIMARY KEY**

- declare the primary key of the relation

Example:

```
PRIMARY KEY (ID)
```

- primary key cannot be null

- **UNIQUE**

- The unique specification states that the attributes  $A_1, A_2, \dots, A_m$  form a candidate key.
- Candidate keys are permitted to be null (in contrast to primary keys).

Example:

```
UNIQUE (course_id, sec_id)
```

# Integrity Constraints

- **CHECK (P)**
  - where P is a predicate

Example:

ensure that semester is one of fall, winter, spring or summer:

```
create table section (  
    course_id varchar (8),  
    sec id_id varchar (8),  
    semester varchar (6),  
    year numeric (4,0),  
    building varchar (15),  
    room_number varchar (7),  
    time slot id varchar (4),  
    primary key (course_id, sec_id, semester, year),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
);
```

# Integrity Constraints - Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
  - Example: If “Biology” is a department name appearing in one of the tuples in the *course* relation, then there exists a tuple in the *department* relation for “Biology”.
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S.

```
create table course (  
    course_id char(5) primary key,  
    title varchar(20),  
    dept_name varchar(20) references department  
)
```

# Cascading Actions in Referential Integrity

- A foreign key with cascade delete/update in a child table means that **if a record in the parent table is deleted/updated, then the corresponding records in the child table will automatically be deleted/updated.**

```
create table course (  
    ...  
    dept_name varchar(20),  
    foreign key (dept_name) references department  
        on delete cascade  
        on update cascade,  
    ...  
)
```

- alternative actions to cascade: **set null, set default**



# Integrity Constraint Violation during Transactions

- Example:

```
create table person (  
    ID char(10),  
    name char(40),  
    mother char(10),  
    father char(10),  
    primary key ID,  
    foreign key father references person,  
    foreign key mother references person)
```

- How to insert a tuple without causing constraint violation ?
  - insert father and mother of a person before inserting person
  - OR, set father and mother to null initially, update after inserting all persons (not possible if father and mother attributes declared to be **not null**)
  - OR *defer constraint checking and use transactions*

# Built-in Data Types in SQL

- **date**: Dates, containing a year, month and date  
Example: **date** '2005-7-27'
- **time**: Time of day in hours minutes and seconds , in hours, minutes and seconds.  
Example: **time** '09:00:30'  
**time** '09:00:30.75'
- **timestamp**: date plus time of day  
Example: **timestamp** '2005-7-27 09:00:30.75'
- **interval**: period of time  
Example: interval '1' day
  - Subtracting a date/time/timestamp value from another gives an interval value
  - Interval values can be added to date/time/timestamp values

# User Defined Types

- **create type** construct in SQL creates user-defined type

**create type** *Dollars* **as numeric (12,2) final**

Using the custom data type:

```
create table department(  
    dept_name varchar (20),  
    building varchar (15),  
    budget Dollars  
)
```

Specify **FINAL** if no subtypes can be created for this type. (default)

Specify **NOT FINAL** if further subtypes can be defined for this type.

# Domains

- **create domain** construct in SQL-92 creates user-defined domain types

**create domain** *person\_name* **char**(20) **not null**

- Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.

**create domain** *degree\_level* **varchar**(10)  
**constraint** *degree\_level\_test*  
**check** (**value in** ('Bachelors', 'Masters', 'Doctorate'));

# Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
  - **blob**: binary large object - object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
  - **clob**: character large object - object is a large collection of character data
  - When a query returns a large object, a pointer is returned rather than the large object itself.

# Indexing

- Indices are data structures used to speed up access to records with specified values for index attributes. (Indices are used to find rows with specific column values quickly.)

**CREATE INDEX** *student\_dept\_name\_index* **ON** *student*(*dept\_name*);

Example scenario:

```
select *  
from student  
where dept_name = "Physics"
```

can be executed faster by using the index to find the required record without looking at all records of *student*

- The users cannot see the indexes, they are just used to speed up searches/queries.
- Primary key is indexed by default in many implementations (MySQL)
- More on indices in Chapter 11

# Indexing

- Index is a data structure which is implemented using
  - B-Trees
  - Hash tables
  - R-Trees
- Advantages of using an index
  - Faster search time by not having to scan entire table
  - Specially when running SELECT queries or JOINS
- Disadvantages
  - Takes up space - larger the number of rows, larger the index size
  - Need to update the index, when rows are added, deleted or updated
- Best practices
  - Indices should only be used if the data in the indexed column is queried frequently
  - Index columns that are being used as foreign keys in other tables
  - Add additional indices based on performance requirements
  - Do not index every column

Sample entry in an index  
("BIO-319", 0x562189)

Thank you!

---



# Truncate and Delete

- Truncate
  - classified as a DDL statement
  - empties the table completely
  - drop and re-create the table, which is much faster than deleting rows one by one
  - does not invoke ON DELETE triggers
  - cannot be rolled back on MySQL
- Delete
  - classified as a DML statement
  - deletes row by row (can specify conditions with where clause)
  - operations are logged individually
  - invokes on delete triggers
  - can be rolled back